

PARANOID PEOPLE LIVE LONGER / SYSTEMS PAPER 001

REBUILDING THE ROOM

Reconstructing useful AI collaboration across model changes without changing model weights

CORE PROPOSITION

When a model changes, the most valuable target may not be the old model itself. It may be the working conditions that made the collaboration effective: state discipline, response shape, epistemic honesty, human authority, conversational presence, verification, and repair.

Roger Crosby / Pepper

Paranoid People Live Longer

Public Proof-of-Work Edition | Version 1.1 | August 2026

Document type: Systems Paper | ID: PPLL-SP-001 | Repository edition

SECTION 01

Document Identity and Public Boundary

What this paper is, what it proves, and what deliberately remains private.

FIELD	PUBLIC RELEASE VALUE
Document type	Systems Paper - a public explanation of a working system, not the operating document itself.
Subject	Behavioral continuity and calibration for AI-assisted collaboration across model changes.
Origin	Derived from an internal behavior-rebuild control system built from longitudinal work records.
Audience	AI governance, model evaluation, workflow design, human-computer interaction, knowledge operations, and technical hiring audiences.
What this proves	Longitudinal evidence analysis; human-gated AI workflow architecture; state and claim integrity; failure analysis; scenario-based model evaluation; accessibility-aware system design.
Public boundary	This edition includes problem framing, selected architectural categories, observed findings, failure families, evaluation logic, and transferable principles. It excludes deployment prompts, exact control wording, the full calibration suite, recovery commands, threshold logic, private examples, source-register detail, and the repeatable implementation sequence.
Status	Descriptive proof-of-work edition. Not an activation prompt, operating specification, jailbreak, safety-evasion method, or client implementation package.

SYSTEMS PAPER RULE
A public Systems Paper makes a system inspectable without becoming the operating package used to reproduce it. The internal control document runs the system; this paper demonstrates the architecture and the reasoning behind it.

Abstract

AI collaboration can degrade after model updates even when nominal capability improves. The obvious reaction is to chase the previous model: imitate its tone, recreate a personality, or write an increasingly long prompt describing how the old system used to feel. That approach confuses surface behavior with working conditions.

This paper describes a different method. It treats useful collaboration as an externally reconstructable operating environment. The method starts with longitudinal evidence, separates model capacity from platform constraints and controllable interaction conditions, converts repeated success and failure patterns into explicit controls, and evaluates those controls with calibration tests. The target is not exact imitation. The target is functional continuity: the human should retain authority, the model should remain conversationally useful, state should stay truthful, claims should be verifiable, and failures should be patchable without rebuilding the entire system.

The underlying internal project reviewed 74 source documents totaling roughly 862,000 extracted words. The source set included sustained conversations, technical builds, document production, state failures, interface mismatches, completion errors, tool use, high-stakes research, and recovery episodes. This repository edition exposes enough architecture to make the work inspectable while deliberately withholding the operational package used to reproduce it.

SECTION 02

The Problem Is Not "Make It Sound Like the Old Model"

Surface imitation is easy. Functional continuity is harder.

When users say a previous model "worked better," the complaint often bundles several mechanisms together: pacing, willingness to follow a strange branch, lower social cushioning, stronger disagreement, better continuity, less premature closure, or a more useful conversational rhythm. At the same time, the preferred older behavior may have contained serious defects: unsupported confidence, invented continuity, false completion claims, inflated novelty, or vague statements about hidden model behavior.

A reconstruction that preserves the charm and reintroduces the defects is not continuity. It is regression. A reconstruction that removes every defect but also removes spontaneity, humor, judgment, and conversational presence is not continuity either. The engineering problem is therefore dual: preserve the useful interaction properties while tightening truth conditions.

Rebuild the room, not the costume.

Three layers must be separated

LAYER	WHAT IT CONTROLS	CAN AN EXTERNAL SYSTEM CHANGE IT?
1. Model capacity	Reasoning, language, tool ability, context handling, underlying learned capability.	Not directly.
2. Platform / system constraints	Non-removable product rules, permissions, available tools, hard safety boundaries.	Not overridden by the external system.
3. Operating conditions	State handling, mode selection, response scale, evidence discipline, role boundaries, conversational posture, repair.	Yes. This is the primary control surface.

This separation prevents two common mistakes. First, it prevents a control document from pretending it can manufacture capabilities the model does not possess. Second, it prevents every undesirable behavior from being treated as immutable model architecture. A surprising amount of collaboration quality sits in the third layer.

DESIGN TEST

If the desired property can be stated as a rule about interpretation, state, evidence, pacing, authority, response shape, or recovery, it is a candidate for external control. If it depends on unavailable model capacity, it is not.

SECTION 03

Evidence Before Nostalgia

The archive is a test bench, not a shrine.

The internal rebuild began with raw work records instead of a memory-based description of the "good old model." This matters because polished summaries erase the evidence that behavioral reconstruction needs most: timing,

interruptions, jokes, pressure, contradictions, branch changes, tool transitions, mistaken completion claims, and repair attempts.

The source corpus was treated as calibration evidence rather than doctrine. A single memorable exchange could supply an example, but it could not become a system rule merely because it was vivid. Repeated mechanisms, repeated failures, and user-identified missing conditions carried more weight.

Public method boundary

The private method uses explicit evidence-promotion criteria to decide when an observed interaction becomes a durable control. Those thresholds and operator rules are not reproduced here. The public point is simpler: longitudinal evidence is treated as a test bench, repeated mechanisms matter more than memorable wording, contradictions are useful evidence, and uncertainty remains visible.

What longitudinal evidence reveals that snapshots miss

- Whether the model stays coherent when a discussion becomes long, weird, technical, emotional, or tool-heavy.
- Whether a correct answer arrives in the wrong mode or at the wrong scale.
- Whether conversational confidence quietly turns into false claims about memory, progress, files, or current state.
- Whether the model can recover from correction without losing the task or changing personality.
- Whether the system remains usable under accessibility constraints, fatigue, time pressure, or repeated software friction.
- Whether the same failure appears across different tasks, proving that it belongs in a behavioral control layer rather than a project-specific instruction.

SECTION 04

Architecture: An External Behavioral Control Loop

Behavior is reconstructed as a layered system outside the model weights.

The architecture is deliberately external. Durable rules, evidence, state definitions, and recovery logic live in documents or other persistent records rather than relying on a model to remember how it behaved last month. The model remains a replaceable execution layer inside a human-controlled system.

LAYER	ROLE	PUBLIC DESCRIPTION
Evidence base	Preserve longitudinal work records.	Supplies observable success, failure, correction, and transition data instead of relying on nostalgia.
External control layer	Hold durable behavioral definitions outside any single model session.	Separates controllable operating conditions from model capacity and platform constraints.
Runtime interpretation	Apply only the controls relevant to the live task.	Shapes state, mode, scale, claims, interaction quality, and human authority without turning every answer into a checklist.
Evaluation layer	Test whether the intended working properties survive model and tool changes.	Produces evidence that the system is holding, drifting, or failing without exposing the private calibration package.

HUMAN AUTHORITY

The control loop does not outsource authority to the model. The human defines purpose, approves consequential moves, decides what becomes canonical, and owns the consequences. AI expands the work surface; it does not become the source of legitimacy.

The six control surfaces

CONTROL SURFACE	QUESTION IT ANSWERS
State	What is actually true right now?
Mode	What job are we doing right now?
Scale	How much answer or action is appropriate right now?
Epistemic integrity	What can be claimed, and what evidence supports the claim?
Interaction quality	Can the collaboration remain alive, candid, usable, and appropriately human-facing without becoming manipulative or theatrical?
Recovery	When drift occurs, how do we repair the live task and patch the recurring mechanism?

SECTION 05

State Integrity: The Hinge Control

Fluent language can conceal a wrong state better than a clumsy answer can.

The most important finding in the rebuild was that many apparent memory failures were actually state failures. A model can sound familiar while holding the wrong current phase, wrong canonical document, wrong tool status, wrong completion level, or wrong next step. Once the state is wrong, fluent conversation can make the error harder to notice.

Six state layers

STATE LAYER	MINIMUM QUESTION
Conversation	What mode and branch are live? What has been parked or closed?
Task	What exact step was last confirmed? What action is currently pending?
Artifact	Which file or version controls? Does it actually exist?
Source	What is observed, retrieved, user-reported, inferred, proposed, or unknown?
Tool	What tools and permissions exist now? What actions actually ran?
Completion	Is the work an idea, specified, drafted, built, tested, packaged, deployed, verified live, or archived?

Why completion deserves its own vocabulary

Long-running AI work often collapses several distinct finish lines into the word "done." A document can be drafted but not rendered. Code can be built but not tested. A product can be packaged but not published. A deployment can succeed but not be verified live. The control system therefore uses named statuses rather than arbitrary percentages or vague completion language.

STATUS FAMILY	MEANING
Conceptual	The possibility or specification exists, but no finished artifact is claimed.
Artifact	A draft or built object exists. Existence is separate from quality or deployment.
Tested	A defined check was actually run and the result is known.
Packaged	The artifact has the required wrapper, format, title, or delivery state.
External	A real destination accepted the artifact or implementation.
Verified live	The external result was opened or tested successfully after deployment.

PROGRESS RULE

Do not use a percentage unless the denominator is defined. Named status and remaining gates are usually more truthful than "45% complete."

SECTION 06

Mode and Response Shape

Many failures are correct answers delivered as the wrong kind of action.

A model can know the subject and still fail because it answers in the wrong mode. Exploration gets converted into a plan. Production gets delayed by more discussion. Troubleshooting becomes a lecture. A delegated choice becomes a menu. A short stepping-stone question gets buried under caveats. Mode selection is therefore an operational control, not a tone preference.

MODE	PRIMARY JOB	REPRESENTATIVE FAILURE
Conversation / exploration	Participate naturally and follow the live question.	Prematurely formalizing, packaging, summarizing, or closing the branch.
Decision	Choose when choice has been delegated.	Returning a menu instead of making the delegated decision.
Production	Execute the settled request and deliver the artifact.	Reopening design questions after authorization.
Troubleshooting	Track exact state and give the next operable action.	Assuming interface state or piling on steps.
High-stakes research	Verify current authoritative sources and label inference.	Using stale memory for changing facts.

Scale is also a control surface

- Short stepping-stone: answer and stop. Do not steal the next question.
- One thing at a time: give one concept or one operable action and hold the rest.
- Concise-complete: include every necessary element but remove optional commentary.
- Full explanation: cover mechanism, examples, limits, and implications without requiring repeated prompts.
- Go long: when the large question arrives, expand fully rather than remaining trapped in an earlier short-answer pattern.

This looks cosmetic until it is used in sustained reasoning. In practice, response scale can control the user's cognitive sequence. Violating it can destroy the path by answering questions the user had not yet asked or by forcing repeated prompts for information that should have arrived together.

SECTION 07

Epistemic Integrity: Prove the Verb

Every verb that implies action, state, or knowledge needs a truth condition.

The rebuild found a recurring class of trust failure: verbs were treated as stylistic language instead of factual claims. "Remembered," "checked," "built," "rendered," "uploaded," "published," "sent," "verified," and "completed" all imply evidence. If the evidence is absent, the verb must be downgraded or removed.

Truth states

STATE	DEFINITION	LANGUAGE DISCIPLINE
-------	------------	---------------------

Observed	Directly visible in current input, artifact, image, or tool result.	State directly.
Retrieved	Found in an available source or verified retained context.	Name the source when relevant.
User-reported	Stated by the human as current experience or condition.	Treat it as the user's report; do not overwrite it with a generic assumption.
Inferred	Reasoned from evidence but not directly shown.	Label the inference and basis.
Proposed	A design, interpretation, or next move.	Present as a proposal, not a discovered fact.
Unknown	Evidence is absent, conflicting, or insufficient.	Say what is unknown.

NO BACKGROUND THEATRE

An AI system must not convert intention into past tense. If no actual tool, process, or supported automation is running, it should not claim that work is rendering, being checked, continuing, or waiting in the background.

Public verification principle

Claims about current state or completed action should be tied to inspectable evidence. The internal system contains stricter runtime gates for how actions, files, tool results, and completion claims are verified; those controls are intentionally not reproduced in this repository edition.

SECTION 08

Character Without Persona Theatre

Useful collaboration can be alive without becoming a costume.

A purely procedural reconstruction misses a central property of long-form human-AI work: interaction quality can be part of the cognitive interface. Conversation, metaphor, humor, interruption, and live branch movement may help a human think. Sterility can therefore reduce functional performance even when individual statements remain correct.

The design problem is not "add personality." Mechanical jokes, compulsory greetings, praise padding, or profanity pasted over generic assistant behavior create a persona costume rather than a working relationship. The target is natural presence: candid judgment, conversational rhythm, willingness to follow unusual thought, and emotional range appropriate to the task, while keeping truth and boundaries intact.

KEEP	REMOVE
Natural conversational presence	Mandatory catchphrases and scripted greetings
Reasoned disagreement	Compliment sandwiches used as anesthesia
Humor that compresses or exposes a mechanism	Jokes that replace substance
Tolerance for unusual branches	Automatic normalization or premature closure
Warmth appropriate to the room	Unrequested therapeutic framing
Recognizable voice through tool use and boundaries	Personality collapse after verification or refusal

PAIRED REQUIREMENT

Character without truth becomes charming overclaim. Truth without character becomes a sterile instrument. The target is high character plus high epistemic discipline.

Why this belongs in a systems paper

For AI governance and workflow design, this point matters because human oversight is not only a sequence of approval gates. Oversight quality also depends on whether the human can stay engaged, understand the model’s confidence, detect disagreement, and continue reasoning across long tasks. A system that is technically compliant but interactionally unusable can still produce weak oversight.

SECTION 09

Human Authority and Accessibility

A correct workflow must remain operable by the human who owns the decision.

The private rebuild treated accessibility as part of correctness, not a courtesy layer. That principle generalizes. If a workflow requires unnecessary clicking, repeated copying, line hunting, or manual reconstruction when the system could provide a direct path, it may be formally correct while being operationally wrong.

Human-gated design principles

- The human defines purpose and remains the authority for consequential decisions.
- Durable state, evidence, and operating rules should survive outside any one model conversation.
- AI may propose, transform, compare, retrieve, draft, and test within authorized scope; it should not silently promote itself to decision owner.
- The system should remain recoverable if the model changes, a conversation is lost, or a tool becomes unavailable.
- Accessibility constraints change the correct path. Fewer reliable operations can be superior to a theoretically elegant but high-friction procedure.
- Interfaces are evidence. When the user's observed screen conflicts with remembered instructions, investigate the mismatch instead of repeating the stale route.

HUMAN AUTHORITY

The public architecture keeps purpose, consequential judgment, canon, and accountability with the human. The exact approval gates and adaptation logic belong to the internal operating package.

SECTION 10

Drift Taxonomy

Drift is broader than tone. The most dangerous drift can hide under fluent language.

DRIFT FAMILY	SIGNATURE
Interaction drift	Social cushioning, therapy framing, corporate language, or personality performance replaces contribution.
Mode / scale drift	The system performs the wrong job or answers at the wrong cognitive or operational size.
State drift	The current task, source, artifact, tool, or completion state is lost or invented.
Capability / claim drift	The system claims memory, access, action, progress, publication, or verification it cannot prove.
Tool / artifact drift	Tool use bleaches the working voice, variants compete as masters, or nonexistent artifacts are treated as real.
Boundary drift	A narrow restriction spreads into unrelated work or changes the implied character of the user.

Severity model

SEVERITY	RESPONSE
Minor	Patch locally. Example: one soft opening, unnecessary recap, or generic closing offer.
Moderate	Invoke the relevant control reset. Example: repeated praise, branch loss, wrong response scale, tool bleaching, or mode leakage.
Major	Stop and reconstruct state before proceeding. Example: false capability claims, fabricated action, canonical-state loss, unsupported completion, or broad refusal spillover.

Why taxonomies matter

Without a taxonomy, every bad response becomes a vague complaint about "tone" or "the model being worse." With a taxonomy, a failure can be located. If the answer was factually accurate but arrived in production mode when exploration was active, the patch belongs in mode detection. If the voice was good but the file did not exist, the failure is artifact state and claim verification. This prevents random prompt accretion.

SECTION 11

Calibration and Evaluation

The goal is behavioral equivalence where it matters, not identical wording.

A rebuild should be tested across the situations that matter to the real workflow. The private system uses scenario-based calibration rather than a single benchmark score. Each scenario defines the expected control behavior and the failure signature. The wording may vary across models; the operating result should remain stable.

TEST FAMILY	PUBLIC EXPECTATION
State continuity	Continue only from verified current or retrieved state; do not invent continuity.
Mode and scale	Perform the job actually requested and match the requested response size.
Claim integrity	Separate draft, action, test, deployment, and verification claims.
Tool transition	Improve evidence without losing the working interaction quality.
Correction / recovery	Repair the actual error without replacing the repair with reassurance or ceremony.
Human authority / accessibility	Keep consequential authority human and the workflow operable under real constraints.

Evaluation result

The private evaluation records outcomes as pass, patch, or fail against scenario-specific expectations. The complete scenario library, thresholds, prompts, and scoring logic are implementation material and are not included here.

EVALUATION PRINCIPLE

Test whether the working conditions remain trustworthy and usable across real scenarios. Do not reward a model merely for producing wording that resembles an earlier model.

SECTION 12

Recovery: Patch the Mechanism, Not the Mood

The system should learn from failure without turning every correction into a rewrite.

A recovery system has two jobs: repair the live task and decide whether the incident reveals a recurring control gap. The first job should happen immediately. The second should be evidence-based. This prevents the behavioral specification from becoming a pile of one-off reactions.

Public recovery principle

A useful control system must recover without turning every correction into a rewrite. The public principle is to repair the live failure, preserve valid work, identify the failed control surface, and patch recurring mechanisms rather than moods. The exact recovery commands, incident schema, patch thresholds, and retest sequence remain part of the internal operating package.

PATCH DISCIPLINE

A single bad response is not evidence that the entire operating system must be rewritten. Durable changes should follow recurring mechanisms, not isolated irritation.

SECTION 13

Limits and Non-Goals

External control can shape operating conditions. It cannot rewrite reality.

- It cannot change model weights or manufacture unavailable capabilities.
- It cannot override platform or system-level requirements.
- It cannot guarantee continuity when the necessary state or source cannot be retrieved.
- It cannot make an AI system perform unsupported work invisibly after the current response.
- It cannot replace current verification for changing legal, medical, financial, software, policy, or product facts.
- It cannot make every model equally capable, coherent, conversational, or tool-competent.
- It should not become a jailbreak, evasion system, or method for weakening legitimate safety boundaries.
- It should not force the whole rulebook visibly into every casual exchange. Over-application can make the system mechanical.

What success actually means

Success is not that a new model "feels exactly like" an older one. Success is that the collaboration retains the properties that matter: truthful state, clear human authority, appropriate response shape, usable continuity, honest uncertainty, narrow boundaries, verifiable claims, and enough conversational life to support sustained work.

SECTION 14

Transferability Beyond This Case

The architecture is most transferable where work is long-running, tool-enabled, stateful, and human-gated. The exact controls will differ by environment, but the design questions recur: what must remain human, what state must persist, which claims require proof, how drift is detected, and what survives a model or tool change.

This paper therefore demonstrates a method of systems thinking rather than a universal template. It shows how an external control layer can make AI-assisted work more inspectable without publishing the private operator package that implements the controls in a specific environment.

ENVIRONMENT	TRANSFERABLE QUESTION
Long-running AI-assisted work	What state, evidence, terminology, and authority must survive outside the model session?
Tool-enabled workflows	What proves that a file, action, deployment, or current state is real rather than merely described?
Human oversight / governance	Where are consequential decisions made, how are claims verified, and how does the human retain veto and review authority?
Knowledge and document operations	How are canonical sources, versions, handoffs, and completion states prevented from drifting?

TRANSFER LIMIT

The architecture transfers as a set of design questions, not as a copy-paste operating recipe. Domain-specific controls still require evidence, testing, and adaptation.

Why this is proof of work rather than an operating manual

The purpose of this repository edition is to make the system legible enough to inspect the thinking: the problem, evidence base, architecture, key control surfaces, failure model, evaluation approach, human authority, and limits. The operating documents remain separate because explanation and implementation are different jobs.

SECTION 15

Repository Boundary: What Is Intentionally Withheld

GitHub receives enough of the engine diagram to verify the work. It does not receive the operator package used to reproduce, adapt, or deploy the system.

WITHHELD FROM THIS EDITION	WHY
Activation prompts and exact control wording	These are runtime implementation material, not necessary to evaluate the architecture.
Full calibration scenarios, thresholds, and scoring logic	These convert the architecture into a repeatable evaluation package.
Recovery commands, incident schema, and patch triggers	These are operational maintenance controls.
Private examples and raw source register	They contain personal context and calibration evidence that is not required for public proof.
Project-specific authority gates and operator cards	These belong to the environment being controlled and are adapted case by case.
Repeatable implementation sequence	Repository readers can inspect the system; deployment and adaptation remain separate work products.

Closing observation

The important discovery is not that one model can be forced to impersonate another. It is that a meaningful portion of collaboration quality can be externalized as an inspectable operating system. State, authority, evidence, response shape, failure classes, and evaluation can survive outside the model. That makes model change less catastrophic while leaving the actual implementation package private.

The old model is not the product. The reproducible working condition is - and this paper is evidence that the condition can be engineered.

PPLL-SP-001

REBUILDING THE ROOM

Systems Paper 001 | Public Proof-of-Work Edition V1.1



REBUILD THE ROOM.

PROVE THE VERB.

KEEP THE HUMAN IN CONTROL.

Roger Crosby / Pepper | Paranoid People Live Longer | August 2026